

Applied Time Series Analysis

Vipul Bhatt

2018-10-08

Contents

Preface	5
1 Introduction to Forecasting	7
1.1 Time Series	7
1.2 Serial Correlation	7
1.3 Testing for Serial Correlation	9
1.4 White Noise Process	10
1.5 Important Elements of Forecasting	10
1.6 Loss Function and Optimal Forecast	12
2 Regression-based Forecasting	15
2.1 Scenario Analysis and Conditional Forecasts	15
2.2 Unconditional Forecasts	15
2.3 Some practical issues	16
2.4 Distributed Lag Regression Models	16
2.5 Application: A Model of Investment Expenditure	17
3 Components of a Time Series	21
3.1 Decomposing a time series	21
3.2 Uses of Decomposition of a time series	21
4 Smoothing Methods	23
4.1 Moving Average Method	23
4.2 Simple Exponential Smoothing	23
4.3 Holt-Winters Smoothing	24
4.4 Holt-Winters Smoothing with Seasonality	24
4.5 Application	24
5 Modeling Trend and Seasonal Components	27
5.1 Trend Estimation	27
5.2 Seasonal Model	29
6 Modeling Cycle	35
6.1 Stationarity and Autocorrelation	35
6.2 Autoregressive (AR) Model	39
6.3 Estimating an AR model	40
6.4 Moving Average (MA) Model	42
6.5 ARMA(p, q)	43

Preface

These lecture notes are prepared for an upper level undergraduate course in time series econometrics. Every fall I teach a course on applied time series analysis at James Madison University. These notes borrow heavily from the teaching material that I have developed over several years of instruction of this course.

One of my main objective is to develop a primer on time series analysis that is more accessible to undergraduate students than standard textbooks available in the market. Most of these textbooks in my opinion are densely written and assume advanced mathematical skills on the part of our students. Further, I have also struggled with their topic selection and organization. Often I end up not following the chapters in order and modify content (by adding or subtracting) to meet my students needs. Such changes causes confusion for some students and more importantly discourages optimal use of the textbook. Hence, this is an undertaking to develop a primer on time series that is accessible, follows a more logical sequencing of topics, and covers content that is most useful for undergraduate students in business and economics.

Note: These notes have been prepared by me using various sources, published and unpublished. All errors that remain are mine.

Chapter 1

Introduction to Forecasting

1.1 Time Series

A time series is a specific kind of data where observations of a variable are recorded over time. For example, the data for the U.S. GDP for the last 30 years is a time series data.

Such data shows how a variable is changing over time. Depending on the variable of interest we can have data measured at different frequencies. Some commonly used frequencies are intra-day, daily, weekly, monthly, quarterly, semi-annual and annual. Figure 1.1 below plots data for quarterly and monthly frequency.

The first panel shows data for the real gross domestic product (GDP) for the US in billions of 2012 dollars, measured at a quarterly frequency. The second panel shows data for the advance retail sales (millions of dollars), measured at monthly frequency.

Formally, we denote a time series variable by y_t , where $t = 0, 1, 2, \dots, T$ is the observation index. For example, at $t = 10$ we get the tenth observation of this time series, y_{10} .

1.2 Serial Correlation

Serial correlation (or auto correlation) refers to the tendency of observations of a time series being correlated over time. It is a measure of the temporal dynamics of a time series and addresses the following question: what is the effect of past realizations of a time series on the current period value? Formally,

$$\rho(s) = Cor(y_t, y_{t-s}) = \frac{Cov(y_t, y_{t-s})}{\sqrt{\sigma_{y_t}^2 \times \sigma_{y_{t-s}}^2}} \quad (1.1)$$

where $Cov(y_t, y_{t-s}) = E(y_t - \mu_{y_t})(y_{t-s} - \mu_{y_{t-s}})$ and $\sigma_{y_t}^2 = E(y_t - \mu_{y_t})^2$

Here, $\rho(s)$ is the serial correlation of order s . For example, $s = 1$ implies *first order* serial correlation between y_t and y_{t-1} , $s = 2$ implies *second order* serial correlation between y_t and y_{t-2} , and so on.

Note that often we use historical data to forecast. If there is no serial correlation, then past can offer no guidance for the present and future. In that sense, presence of serial correlation of some order is the first condition for being able to forecast a time series using its historical realizations.

Now, we can either have positive or negative serial correlation in data. Figure ?? plots two time series with positive and negative serial correlation, respectively.

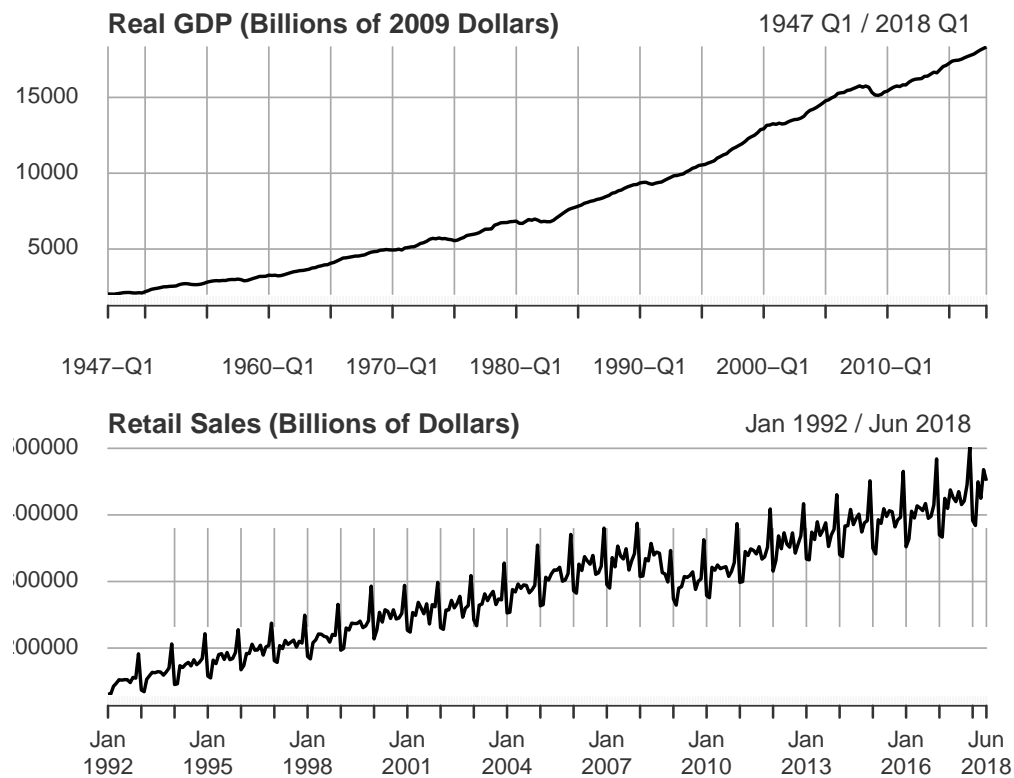


Figure 1.1: Time Series at quarterly and monthly frequency




Figure 1.2: Serial Correlation




Figure 1.3: Serial Correlation

1.3 Testing for Serial Correlion

We can use a Lagrange-Multiplier (LM) test for detecting serial correlation. This test is also known as *Breuch-Godfrey* test. I will use the linear regression model to explain this test. Consider the following regression model:

$$y_t = \beta_0 + \beta_1 X_{1t} + \epsilon_t \quad (1.2)$$

Consider the following model for serial correlation of order p for the error term:

$$\epsilon_t = \rho_1 \epsilon_{t-1} + \rho_2 \epsilon_{t-2} + \dots + \rho_p \epsilon_{t-p} + \nu_t \quad (1.3)$$

Then we are interested in the following test:

$$H_0 = \rho_1 = \rho_2 = \dots = \rho_p = 0$$

$$H_A = \text{Not } H_0$$

To implement this test, we estimate the BG regression model given by:

$$e_t = \alpha_0 + \alpha_1 X_{1t} + \rho_1 e_{t-1} + \rho_2 e_{t-2} + \dots + \rho_p e_{t-p} + \nu_t \quad (1.4)$$

where we replacr the error term with the OLS residuals (denoted by e). The LM test statistic is given by:

$$LM = N \times R_{BG}^2 \sim \chi_p^2$$

If the test statistic value is greater than the critical value then we reject the null hypothesis.

1.4 White Noise Process

A time series is a *white noise* process if it has zero mean, constant and finite variance, and is serially uncorrelated. Formally, y_t is a white noise process if:

1. $E(y_t) = 0$
2. $Var(y_t) = \sigma_y^2$
3. $Cov(y_t, y_{t-s}) = 0 \forall s \neq t$

We can compress the above definition as: $y_t \sim WN(0, \sigma_y^2)$. Often we assume that the unexplained part of a time series follows a white noise process. Formally,

$$Time\ Series = Explained + White\ Noise \quad (1.5)$$

By definition we cannot forecast a white noise process. An important diagnostics of model adequacy is to test whether the estimated residuals are white noise (more on this later).

1.5 Important Elements of Forecasting

Definition 1.1 (Forecast).

A *forecast* is an *informed* guess about the unknown future value of a time series of interest. For example, what is the stock price of Facebook next Monday?

There are three possible types of forecasts:

1. *Density Forecast*: we forecast the entire probability distribution of the possible future value of the time series of interest. Hence,

$$F(a) = P[y_{t+1} \leq a] \quad (1.6)$$

give us the probability that the 1-period ahead future value of y_{t+1} will be less than or equal to a . For example, the future real GDP growth could be normally distributed with a mean of 1.3% and a standard deviation of 1.83%. Figure 1.4 below plots the density forecast for real GDP growth.

2. *Point Forecast*: our forecast at each horizon is a single number. Often we use the expected value or mean as the point forecast. For example, the point forecast for the 1-period ahead real GDP growth can be the mean of the probability distribution of the future real GDP growth:

$$f_{t,1} = 1.3 \quad (1.7)$$

3. *Interval Forecast*: our forecast at each horizon is a range which is obtained by adding *margin of errors* to the point forecast. With some probability we expect our future value to fall within this range. For example, the 95% interval forecast for the next period real GDP growth is (-2.36%, 4.96%). Hence, with 95% confidence we expect next period GDP to fall between -2.36% and 4.96%.

Definition 1.2 (Forecast Horizon).

Forecast Horizon is the number of periods into the future for which we forecast a time series. We will denote it by h . Hence, for $h = 1$, we are looking at 1-period ahead forecast, for $h = 2$ we are looking at 2-period ahead forecast and so on.

Formally, for a given time series y_t , the h -period ahead unknown value is denoted by y_{t+h} . The forecast of this value is denoted $f_{t,h}$.




Figure 1.4: Density Forecast for Future Real GDP Growth




Figure 1.5: Forecast Horizon

Definition 1.3 (Forecast Error).

A *forecast error* is the difference between the realization of the future value and the previously made forecast. Formally, the h -period ahead forecast error is given by:

$$e_{t,h} = y_{t+h} - f_{t,h} \quad (1.8)$$

Hence, for every horizon, we will have a forecast and a corresponding forecast error. These errors can be negative (indicating over prediction) or positive (indicating under prediction).

Definition 1.4 (Information Set).

Forecasts are based on *information* available at the time of making the forecast. *Information Set* contains all the relevant information about the time series we would like to forecast. We denote the set of information available at time T by Ω_T . There are two types of information sets:

1. Univariate Information set: Only includes historical data on the time series of interest:

$$\Omega_T = \{y_T, y_{T-1}, y_{T-2}, \dots, y_1\} \quad (1.9)$$

2. Multivariate Information set: Includes historical data on the time series of interest as well as any other variable(s) of interest. For example, suppose we have one more variable x that is relevant for forecasting y . Then:

$$\Omega_T = \{y_T, x_T, y_{T-1}, x_{T-1}, y_{T-2}, x_{T-2}, \dots, y_1, x_1\} \quad (1.10)$$

1.6 Loss Function and Optimal Forecast

Think of a forecast as a solution to an *optimization* problem. When forecasts are wrong, the person making the forecast will suffer some *loss*. This loss will be a function of the magnitude as well as the sign of the *forecast error*. Hence, we can think of an *optimal forecast* as a solution to a minimization problem where the forecaster is minimizing the loss from the forecast error.

Definition 1.5 (Loss Function).

A *loss* function is a mapping between forecast errors and their associated losses. Formally, we denote the h -period ahead loss function by $L(e_{t,h})$. For a function to be used as a loss function, three properties must be satisfied:

1. $L(0) = 0$
2. $\frac{dL}{de} > 0$
3. $L(e)$ is a continuous function.

Two types of loss functions are:

- Symmetric Loss Function: both positive and negative forecast errors lead to same loss. See Figure 1.6. A commonly used loss function is *quadratic loss function* given by:

$$L(e_{t,h}) = e_{t,h}^2 = (y_{t+h} - f_{t,h})^2 \quad (1.11)$$

- Asymmetric Loss Function: loss depends on the sign of the forecast error. For example, it could be that positive errors produce greater loss when compared to negative errors. See the function below and Figure 1.7 that attaches a higher loss to positive errors:




Figure 1.6: Quadratic Loss Functions




Figure 1.7: Asymmetric Loss Function

$$L(e_{t,h}) = e_{t,h}^2 + 4 \times e_{t,h} \quad (1.12)$$

Once we have chosen our loss function, the optimal forecast can be obtained by minimizing the expected loss function.

Definition 1.6 (Optimal Forecast).

An *optimal forecast* minimizes the expected loss from the forecast, given the information available at the time. Mathematically, we denote it by $f_{t,h}^*$ and it solves the following minimization problem:

$$\min_{f_{t,h}} E(L(e_{t,h})|\Omega_t) \quad (1.13)$$

In theory we can assume any functional form for the loss function and that will lead to a different *optimal forecast*. An important result that follows from a specific functional form is stated as Theorem 1.1.

Theorem 1.1. *If the loss function is quadratic then the optimal forecast is the conditional mean of the time series of interest. Formally, if $L(e_{t,h}) = e_{t,h}^2$ then,*

$$f_{t,h}^* = E(y_{t+h}|\Omega_t) \quad (1.14)$$

Note that $E(e_{t,h}^2)$ is known as *mean squared errors (MSE)*. Hence, the expected loss from a quadratic loss function is the same as the MSE. In this course, we assume that the forecaster faces a quadratic loss function and hence based on Theorem 1.1, we will learn different models for estimating the conditional mean of the future value of the time series of interest, i.e., $E(y_{t+h}|\Omega_t)$.

Chapter 2

Regression-based Forecasting

One way to compute the conditional expectation is the linear regression model. Here, our information set contains data on all relevant explanatory variables available at the time of forecast, i.e.,

$$\Omega_t = X_{1t}, X_{2t}, \dots, X_{Kt} \quad (2.1)$$

Hence, we get the following equality:

$$E(y_t|\Omega_t) = E(y_t|X_{1t}, X_{2t}, X_{3t}, \dots, X_{Kt}) \quad (2.2)$$

The right hand side of the above equation is the multiple regression model of the form:

$$y_t = \beta_0 + \beta_1 X_{1t} + \beta_2 X_{2t} + \dots + \beta_K X_{Kt} + \epsilon_t \quad (2.3)$$

We can easily estimate the above model using Ordinary Least Squares (OLS) and compute the *predicted value* of y :

$$\hat{y}_t = \hat{\beta}_0 + \hat{\beta}_1 X_{1t} + \hat{\beta}_2 X_{2t} + \dots + \hat{\beta}_K X_{Kt} \quad (2.4)$$

The above equation can be used to compute the optimal forecast. Suppose, we are interested in computed the h period ahead forecast for y . Then, using the above equation we get:

$$\hat{y}_{t+h} = \hat{\beta}_0 + \hat{\beta}_1 X_{1t+h} + \hat{\beta}_2 X_{2t+h} + \dots + \hat{\beta}_K X_{Kt+h} \quad (2.5)$$

2.1 Scenario Analysis and Conditional Forecasts

One way to use a regression model to produce forecasts is called *scenario analysis* where we produce a different forecast for the dependent under each possible scenario about the future values of the independent variables. For example, what will be the forecast for inflation if the Federal Reserve Bank raises the interest rate? Would our forecast differ depending on the size of the increase in the interest rate?

2.2 Unconditional Forecasts

An alternative is to separately forecast each independent variable and then compute the forecast for the dependent variable. Yet another alternative is to use lagged variables as independent variables. Depending on the number of lags, we can forecast that much ahead into future (see Distributed Lag Section for details).

2.3 Some practical issues

1. To forecast the dependent variable we first need to compute a forecast for the independent variable. Errors in this step induce errors later.
2. *Spurious regression*: It is quite possible to find a strong linear relationship between two completely unrelated variables over time if they share a common time trend.
3. *Model Uncertainty*: We do not know the true functional form for the regression model and hence our estimated model is only a proxy for the true model.
4. *Parameter Uncertainty*: This kind of forecast uses regression coefficients that are computed using a fixed sample. Over time with new data, there will be changes in these coefficients.

2.4 Distributed Lag Regression Models

Consider the following simple regression model:

$$y_t = \beta_0 + \beta_1 x_t + \epsilon_t \quad (2.6)$$

Here, if want to forecast y_{t+1} then we must either consider different scenarios for x_{t+1} or independently forecast x_{t+1} first, and then use it to compute forecast for y_{t+1} . An alternative is to estimate the following lagged regression model:

$$y_t = \beta_0 + \beta_1 x_{t-1} + \epsilon_t \quad (2.7)$$

Note that by estimating the above model we get the following predicted value equation for $t + 1$:

$$\widehat{y_{t+1}} = \widehat{\beta_0} + \widehat{\beta_1} x_t \quad (2.8)$$

Hence, we can easily produce 1-period ahead forecast from this model. In order to produce forecast farther into future we would need to add more lags of the independent variable to the model. A generalized model of this kind is called *distributed lag model* and is given by:

$$y_t = \beta_0 + \sum_{s=1}^p \beta_s x_{t-s} + \epsilon_t \quad (2.9)$$

The number of lags to include can be determined using some kind of goodness of fit measure.

2.4.1 Dynamic Effect of X on Y

A very useful benefit of estimating a distributed lag model is that it allows us to measure how changes in x in the current period can impact the dependent variable over time. Consider a simple distributed lag model with two lags:

$$y_t = \beta_0 + \beta_1 x_{t-1} + \beta_2 x_{t-2} + \epsilon_t \quad (2.10)$$

In this model the lag structure implies that any change in x will persist for two periods in terms of its effect on y . In fact we now have to consider the *dynamic* effect of x on y . Formally, there are two types of effects:

1. *dynamic effect* of x on y given by:

$$\frac{\partial y_{t+s}}{\partial x_t} \quad s = 0, 1, 2, \dots$$

In our example, the sequence of dynamic effects are:

$$\frac{\partial y_t}{\partial x_t} = 0; \frac{\partial y_{t+1}}{\partial x_t} = \beta_1; \frac{\partial y_{t+2}}{\partial x_t} = \beta_2; \frac{\partial y_{t+s}}{\partial x_t} = 0 \quad \forall s > 2 \quad (2.11)$$

2. *long run effect* of x on y given by:

$$\sum_{s=0}^p \frac{\partial y_{t+s}}{\partial x_t} \quad (2.12)$$

In our example, the long run effect is:

$$\beta_1 + \beta_2$$

2.4.2 Model Selection Criterion

Most often we compare models that have different number of independent variables. For example, in our application, in order to select the number of lags for output and capital stock, we will essentially compare models with different number of independent variables. In such cases we must account for the trade-off between goodness of fit and degrees of freedom. Increasing the number of independent variables will:

1. lower the MSE and hence leads to better fit.
2. lowers the degrees of freedom

Two commonly used measures based on MSE incorporate this trade-off:

1. Akaike Information Criterion (AIC):

$$AIC = MSE \times e^{\frac{2k}{T}}$$

where k is the number of estimated parameters, T is the sample size. Then, k/T is the number of parameters estimated per observation and $e^{\frac{2k}{T}}$ is the *penalty factor* imposed on adding more variables to the model. As we increase k , this penalty factor will increase exponentially for a given value of T .

2. Bayesian Information Criterion (BIC):

$$BIC = MSE \times T^{\frac{k}{T}}$$

Lower values of either AIC or BIC indicates greater accuracy. So we select a model with lower value of either of these two criteria. Note that the penalty imposed by BIC is harsher and hence it will typically select a more parsimonious model (Figure 2.1).

2.5 Application: A Model of Investment Expenditure

2.5.1 A Multiple Regression Model of Investment Expenditure

Suppose have annual data on private investment, private sector output, and capital stock. Our model specification is given by:

$$y_t = \beta_0 + \beta_1 x_{1t} + \beta_2 x_{2t} + \epsilon_t \quad (2.13)$$

We can estimate the above model using OLS and then conduct scenario-based forecasting. For ease of interpretation, we will convert all variables in natural logarithms.

Table 2.1 below presents the estimated coefficients of our regression model. Higher output and capital stock leads to greater investment expenditure.

Next, we forecast of investment expenditure under three different scenarios:




Figure 2.1: Penalty Factor of AIC and BIC

Table 2.1: A Multiple Regression Model of Investment Expenditure

	Estimated Coefficients	Std. Error	t-ratio	p-value
(Intercept)	-4.8421855	0.9623332	-5.031714	0.0000044
x1	0.9987751	0.2418282	4.130102	0.0001104
x2	0.4204833	0.3643054	1.154205	0.2528456

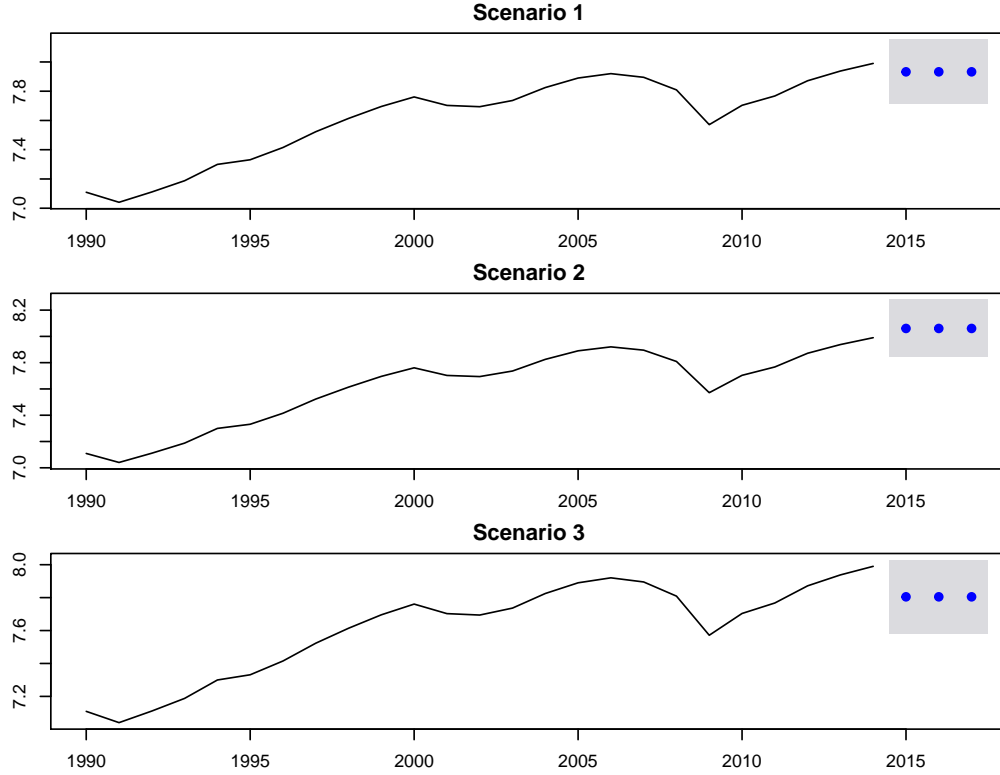


Figure 2.2: Investment outlook for next 3 years

1. For next 3 years, both output and capital stock remain at the average of last 3 years.
2. For next 3 years, both output and capital stock remain at 1% above the average of last 3 years.
3. For next 3 years, both output and capital stock remain at 1% below the average of last 3 years.

Figure 2.2 below present our investment expenditure outlook under these 3 scenarios.

2.5.2 A Distributed Lag Model of Investment Expenditure

In this application we will estimate a distributed lag model for investment expenditure. The idea here is that it takes time for investment to respond to output and capital stock changes. The model specification we want to estimate is:

$$y_t = \beta_0 + \sum_{i=1}^p \beta_i x_{1t-i} + \sum_{i=1}^p \alpha_i x_{2t-i} + \epsilon_t \quad (2.14)$$

where y denotes real investment expenditure of the private sector, x_1 denotes output of the private sector, and x_2 denotes capital stock of the private sector.

We estimate our model by first selecting the optimal lag order for each independent variable, and selecting the one with lowest value for AIC/BIC. From @??tab:ch2-table2) we find that the lowest BIC occurs at lag=2. Hence, we estimate a model with two lags for each independent variable in our model.

Hence, our final model is given by:

$$y_t = \beta_0 + \sum_{i=1}^2 \beta_i x_{1t-i} + \sum_{i=1}^2 \alpha_i x_{2t-i} \quad (2.15)$$

Table 2.2: Optimal Order of the lags

Lag	AIC	BIC
1	-98.42469	-89.72714
2	-127.36043	-114.31411
3	-129.95355	-112.55845
4	-127.49885	-105.75498

Table 2.3: Distributed Lag Model of Investment Expenditure

	Estimated Coefficients	Std. Error	t-ratio	p-value
(Intercept)	-6.541336	0.8300877	-7.880295	0.0000000
L(x1, 1:2)1	10.932987	1.9354346	5.648854	0.0000005
L(x1, 1:2)2	-9.769632	1.9286226	-5.065601	0.0000044
L(x2, 1:2)1	2.545657	0.7735580	3.290842	0.0017028
L(x2, 1:2)2	-2.189088	0.7904689	-2.769354	0.0075309

The results of our estimation are presented below in Table 2.3

1. Using our estimated model we can easily compute the dynamic effect as well as the long run effect of each independent variable on the dependent variable.
2. Given the lag structure of our estimated model, we can also produce forecasts for y_{t+1} by computing the following equation:

$$f_{t,1} = \widehat{y_{t+1}} = \hat{\beta}_0 + \hat{\beta}_1 x_{1t} + \hat{\beta}_2 x_{1t-1} + \hat{\alpha}_1 x_{2t} + \hat{\alpha}_2 x_{2t-1} \quad (2.16)$$

Chapter 3

Components of a Time Series

A given time series can have four possible components:

1. Trend: denoted by B_t captures the long run behavior of the time series of interest.
2. Season: denoted by S_t are *periodic* fluctuations over *seasons*. The period of the season is fixed and known. For example, rise in non-durable sales during Christmas.
3. Cycle: denoted by C_t are *non-periodic* are fluctuations in that they occur regularly but over periods that are not fixed in duration.
4. Irregular: denoted by ϵ_t are random fluctuations, typically modeled as a white noise process.

3.1 Decomposing a time series

We can decompose any given time series into its components. There are two ways to accomplish this:

1. Additive Decomposition: Here it is assumed that all four components are added to obtain the underlying timer series:

$$y_t = B_t + S_t + C_t + \epsilon_t \quad (3.1)$$

2. Multiplicative Decomposition: Here it is assumed that all four components are multiplied to obtain the underlying timer series:

$$y_t = B_t \times S_t \times C_t \times \epsilon_t \quad (3.2)$$

Note that using properties of logarithms, multiplicative decomposition is the same as additive decomposition in log terms:

$$\log(y_t) = \log(B_t) + \log(S_t) + \log(C_t) + \log(\epsilon_t) \quad (3.3)$$

Most statistical software can implement these decomposition using data on a time series variable as input. Typically they combine cyclical component with irregular component and provide a three-way decomposition. In Figure 3.1 I use R to decompose real GDP for the US into its components.

3.2 Uses of Decomposition of a time series

The usefulness of decomposing a time series depends on our objective.




Figure 3.1: Additive Decomposition of Retail Sales

1. It may be of interest to study each component separately or to simply improve our understanding of the temporal dynamics of a time series of interest. Decomposing it into different components is the first step towards achieving that goal.
2. We can also use the decomposition to filter out components that we are not interested in studying. If for example we are only interested in modeling the cyclical component of the time series, then we can assume some kind decomposition, additive or multiplicative, and filter out the trend and seasonal component. For example, assuming additive decomposition, the filtered time series is given by:

$$\text{Filtered } y_t = y_t - B_t - S_t \quad (3.4)$$

We can then proceed to model the cyclical component using the filtered data.

Chapter 4

Smoothing Methods

One way to approach forecasting is to *average* out the fluctuations in the underlying time series to produce a *smoothed* data which can be extrapolated to produce forecasts. These smoothing methods are essentially *model-free* and may not even produce *optimal forecasts*. Depending on the method used one can accommodate seasonal as well as trend components of the underlying time series.

4.1 Moving Average Method

We compute an average of most recent data values for the time series and use it as a forecast for the next period.

An important parameter is the *window* over which we take the average. Let us denote this window by m , then:

$$y_{t+1}^s = \frac{\sum_{i=t-m+1}^t y_i}{m} \quad (4.1)$$

A larger value of m produces greater smoothing and most software have a default value of this parameter which can be changed if needed.

4.2 Simple Exponential Smoothing

In the moving average method, all observations received same weight. However, it is reasonable to argue that more recent observations may have a greater influence than those in the remote past. In this method, the weight attached to past observations exponentially decay over time. Here is the algorithm for computing the smoothed data and its forecast:

1. Initialize at $t=1$:

$$y_1^s = y_1$$

2. Update:

$$y_t^s = \alpha y_t + (1 - \alpha)y_{t-1}^s \quad \text{for } t = 2, 3, \dots, T$$

3. h -period ahead forecast:

$$f_{T,h} = y_T^s$$

Here the h -period ahead forecast is:

Exercise: Can you show that y_t^s is a weighted moving average of all past observations? Use backward substitution method.

Here $\alpha \in (0, 1)$ is the smoothing parameter, with smaller value indicating greater smoothing.

4.3 Holt-Winters Smoothing

We add trend component to the simple exponential smoothing. In step 2 the equation we use to update the smoothed data is given by:

$$\begin{aligned} y_t^s &= \alpha y_t + (1 - \alpha)(y_{t-1}^s + B_{t-1}) \\ B_t &= \beta(y_t^s - y_{t-1}^s) + (1 - \beta)B_{t-1} \end{aligned} \quad (4.2)$$

We now have an additional parameter β that is the trend parameter. Here the h -period ahead forecast is:

$$f_{T,h} = y_T^s + h \times B_T \quad (4.3)$$

4.4 Holt-Winters Smoothing with Seasonality

We now add seasonal component along with trend. Assuming multiplicative seasonality with period n :

$$y_t^s = \alpha \frac{y_t}{S_{t-n}} + (1 - \alpha)(y_{t-1}^s + B_{t-1}) \quad (4.4)$$

$$B_t = \beta(y_t^s - y_{t-1}^s) + (1 - \beta)B_{t-1} \quad (4.5)$$

$$S_t = \gamma \frac{y_t}{y_t^s} + (1 - \gamma)S_{t-n} \quad (4.6)$$

The h -period ahead forecast is given by:

$$f_{T,h} = (y_T^s + h \times B_T) \times S_{T+h-n} \quad (4.7)$$

4.5 Application

We use R to implement a 12-period ahead forecast for new housing starts for the U.S. The data is at monthly frequency from June 2000 through June 2018. The resulting forecasts are plotted in Figure 4.1.

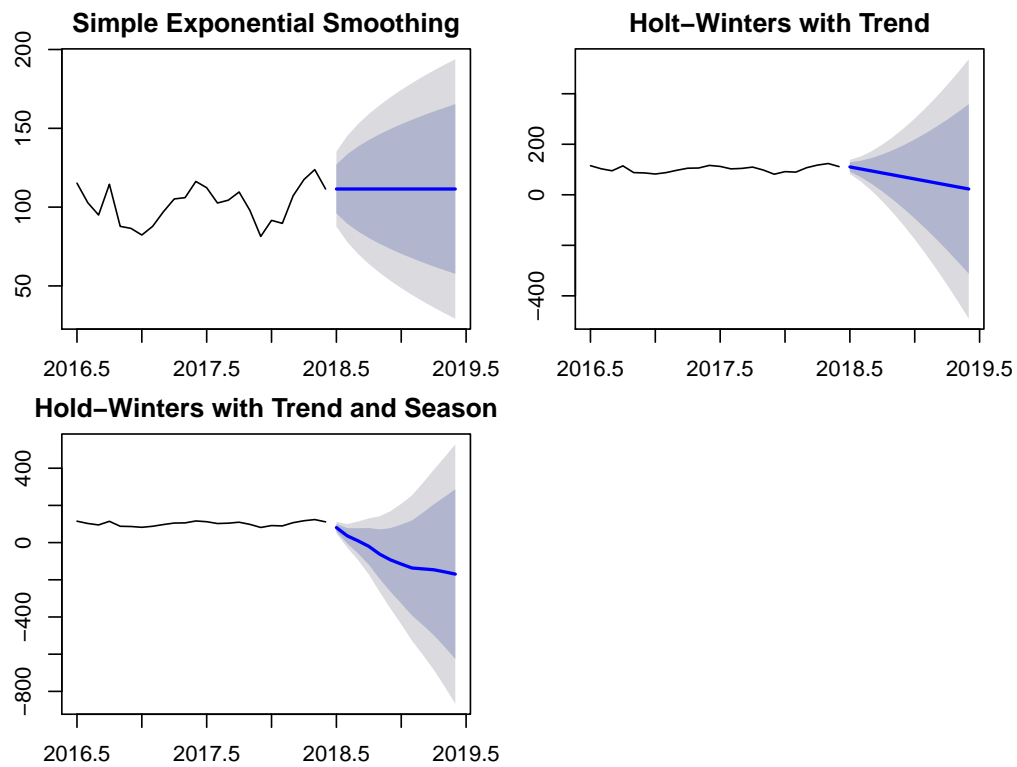


Figure 4.1: Forecast of Housing Starts: Three Smoothing Methods

Chapter 5

Modeling Trend and Seasonal Components

5.1 Trend Estimation

An important component of a time series is *trend* that captures the long run evolution of the variable of interest. There are two types of trends:

1. Deterministic Trend: the underlying trend component is a *known* function of time with *unknown* parameters.
2. Stochastic Trend: the trend component is random.

In this note we will focus on estimating and forecasting deterministic trend models. We will come back to stochastic trend later when we talk about stationarity property of a time series.

5.1.1 Parametrizing a deterministic trend

Whether or not there is deterministic trend in the data can be typically gleaned by simply plotting the time series over time. For example, Figure @ref(fig: ch5-figure1) below plots real GDP for the US at quarterly frequency. We can observe a positive time trend with real GDP increasing with time. In this section we will learn to *fit* a function that captures this relationship accurately.

Note: The variable time is denoted by t and it is artificially created to take value of 1 for the first period, 2 for the second period and so on.

There are two commonly used functional forms for capturing a deterministic trend:

1. Polynomial Trend: We fit a polynomial of appropriate order to capture the time trend. For example,
 - A. Linear trend:

$$y_t = \beta_0 + \beta_1 t + \epsilon_t \quad (5.1)$$

- B. Quadratic trend:

$$y_t = \beta_0 + \beta_1 t + \beta_2 t^2 + \epsilon_t \quad (5.2)$$

In general, we can fit a polynomial of order q :

$$y_t = \beta_0 + \sum_{i=1}^q \beta_i t^i + \epsilon_t \quad (5.3)$$




Figure 5.1: Real GDP (2012 Chained Billions of Dollars)

We can estimate this model using the OLS. One of the key component here is to determine the *right* order of the polynomial. We can begin with a large enough number for q and then select the appropriate order using AIC or BIC criterion.

2. Exponential or log-linear trend: In some cases we may want to use an exponential trend or equivalently a log-linear trend.

$$y_t = e^{(\beta_0 + \beta_1 t + \epsilon_t)} \quad (5.4)$$

$$\text{equivalently} \quad (5.5)$$

$$\log(y_t) = \beta_0 + \beta_1 t + \epsilon_t \quad (5.6)$$

Again we can estimate the above model using OLS.

5.1.2 Uses of the Deterministic Trend Model

Once we have finalized our deterministic trend model i.e., either a polynomial of a specific order or log-linear trend, we can use the estimated model for the following two purposes:

1. Detrending our data: Suppose we would like to eliminate trend from our data. The residual from our final trend model is the *detrended* time series.
2. Forecasting: We can also forecast our time series based on the estimated trend. For example, suppose our final model is a quadratic trend. The predicted value is given by:

$$\hat{y}_t = \hat{\beta}_0 + \hat{\beta}_1 t + \hat{\beta}_2 t^2 \quad (5.7)$$

Then, the 1-period ahead forecast for y_{t+1} can be obtained by solving:

$$\hat{y}_{t+1} = \hat{\beta}_0 + \hat{\beta}_1(t+1) + \hat{\beta}_2(t+1)^2 \quad (5.8)$$

Table 5.1: Optimal Order of the Polynomial

order	AIC	BIC
1	4716.836	4727.794
2	4133.469	4148.079
3	4105.087	4123.350
4	4002.212	4024.127

Table 5.2: Regression Results

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1714.277	81.019	21.159	0
trend	43.398	3.911	11.097	0
I(trend^2)	-0.344	0.055	-6.194	0
I(trend^3)	0.003	0.000	10.233	0
I(trend^4)	0.000	0.000	-11.160	0

5.1.3 Application: Estimating a polynomial trend for U.S. Real GDP

We will now fit a polynomial trend to the US real GDP data that was presented in Figure ???. We first estimate polynomials of different orders and select the optimal order determined by the lowest possible AIC/BIC. Table 5.1. shows these statistics for up to 4th order polynomial. We find that the lowest value occur at $q = 4$.

Hence, our final trend model is:

$$y_t = \beta_0 + \beta_1 t + \beta_2 t^2 + \beta_3 t^3 + \beta_4 t^4 + \epsilon_t \quad (5.9)$$

The estimated trend model is presented in Table 5.2.

Using the estimated model, we can compute the detrended data as the residual and also forecast y_t . Figure 5.2 below plots the detrended real GDP obtained as a residual from our trend model.

Figure 5.3 shows the forecast of real GDP for next 8 quarters along with the 95% confidence bands.

5.2 Seasonal Model

We now focus on the *seasonal* component of a time series, i.e., that is periodic fluctuations that repeat themselves every season. For example, increase in ice cream sales during summer season. Just like trend component, such seasonal pattern could be *deterministic* or *stochastic*. In this chapter we will focus on estimating deterministic seasonal component.

In Figure 5.4 we plot housing starts in the U.S. The data is at monthly frequency and we can see a clear seasonal pattern. Housing starts seem to increase in spring and summer months. This is followed by a decline in fall and winter months.

One option to deal with seasonality is to either obtain seasonally adjusted data from the source itself. Alternatively, we can use decomposition method and appropriately filter out the seasonal component. However, if our objective is to explicitly model the seasonal component of a time series then we must work with non-seasonally adjusted data.




Figure 5.2: Detrended Real GDP




Figure 5.3: Forecast of Real GDP




Figure 5.4: Housing Starts in U.S.

5.2.1 Regression Model with Seasonal Dummy Variables

One way to account for seasonal patterns in data is to add dummy variables for season. To avoid perfect multicollinearity, if there are s seasons, we can include $s - 1$ dummy variables. For example, for quarterly data, $s = 4$ and hence we need $s - 1 = 3$ dummy variables in our regression model. Formally, for quarterly data, the seasonal regression model is given by:

$$y_t = \beta_0 + \beta_1 D_{1t} + \beta_2 D_{2t} + \beta_3 D_{3t} + \epsilon_t \quad (5.10)$$

In the above regression model, D_1 , D_2 , and D_3 are dummy variables that capture first three quarters of the year. For example, $D_1 = 1$ for the first quarter and $D_1 = 0$ otherwise. Similarly, $D_2 = 1$ for the second quarter and $D_2 = 0$ otherwise. In this example, we use the fourth quarter as the *base group*.

The above model can be estimated using OLS. Again, we can use the residual from our estimated model as a measure of *deseasonalized* data. We can also forecast the dependent variable based on the seasonal component only.

5.2.2 Application: Seasonal Model of Housing Starts

We now estimate a seasonal regression model for the housing starts data presented in Figure 5.4. The data is at monthly frequency which implies we can have 12 possible seasons and hence would need 11 dummy variables in our regression model. Formally, we use January as the base group and include dummy variables for the last 11 months of the year:

Table 5.3: Regression Results

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	84.422	9.852	8.569	0.000
season2	2.283	13.933	0.164	0.870
season3	19.133	13.933	1.373	0.171
season4	28.350	13.933	2.035	0.043
season5	34.061	13.933	2.445	0.015
season6	35.562	13.749	2.587	0.010
season7	33.422	13.933	2.399	0.017
season8	27.794	13.933	1.995	0.047
season9	26.139	13.933	1.876	0.062
season10	25.622	13.933	1.839	0.067
season11	10.489	13.933	0.753	0.452
season12	0.639	13.933	0.046	0.963

$$y_t = \beta_0 + \sum_{i=2}^{12} \beta_i D_{it} + \epsilon_t \quad (5.11)$$

Table 5.3 presents the estimation results for this exercise. In Figure 5.5 we plot the forecast of housing starts for next 12 months using our estimated model, along with 95% confidence bands.




Figure 5.5: Forecast of Housing Starts

Chapter 6

Modeling Cycle

In this chapter we will focus on the cyclical component of a time series and hence focus on data that either has no trend and seasonal components, or data that is filtered to eliminate any trend and seasonality. One of the most commonly used method to model cyclical component is the *Autoregressive Moving Average (ARMA)*. This model has two distinct components:

1. *Autoregressive (AR) component*: the current period value of a time series variable depends on its past (lagged) observations. We use p to denote the **order** of the AR component and is the number of lags of a variable that directly affect the current period value. For example, a firm's production in the current period maybe impacted by past levels of production. If last year's production exceeded demand, the stock of unsold goods may be used to meet this period demand first, hence lowering the current period production.
2. *Moving average (MA) component*: the current period value of a time series variable depends on current period **shock** as well as past shocks to this variable. We use q to denote the **order** of the MA component and is the number of past period shocks that affect the current period value of the variable of interest. For example, if the Federal Reserve Bank raises the interest in 2016, the effects of that policy shock may impact investment and consumption spending in 2017.

Before we consider these time series model in details it is useful to discuss certain properties of time series that allow us a better understanding of these models.

6.1 Stationarity and Autocorrelation

6.1.1 Covariance Stationary Time Series

Definition 6.1 (Covariance Stationary Time Series).

A time series $\{y_t\}$ is said to be a *covariance stationary process* if:

1. $E(y_t) = \mu_y \quad \forall \quad t$
2. $Var(y_t) = \sigma_y^2 \quad \forall \quad t$
3. $Cov(y_t, y_{t-s}) = \gamma(s) \quad \forall \quad s \neq t$

One way to think about stationarity is *mean-reversion*, i.e, the tendency of a time series to return to its *long-run* unconditional mean following a shock (or a series of shock). Figure @??fig:ch6-figure1) below shows this property graphically.

In practice however, you will not be able to visualize a mean-reverting stationary process this clearly. For example, in Figure 6.2 we plot real GDP growth for the U.S. which is a stationary process with a mean of

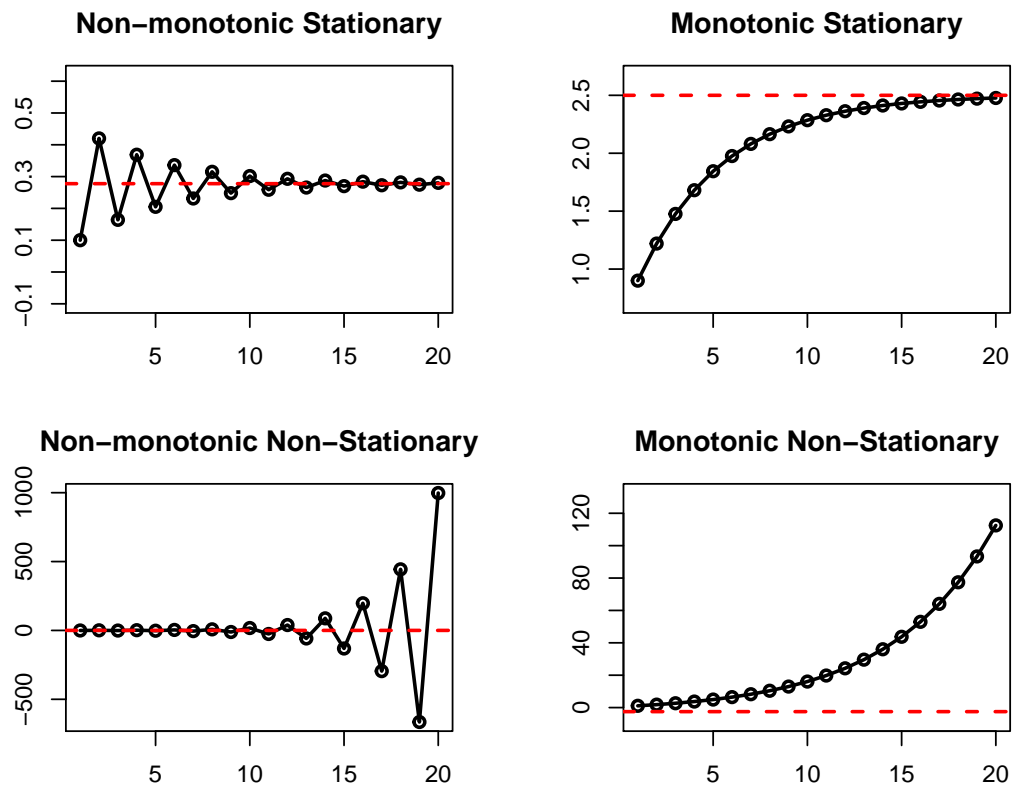


Figure 6.1: Reversion to mean




Figure 6.2: Reversion to mean in practice




Figure 6.3: ACF for a Stationary Time Series

0.7%. In this chapter we will only consider stationary time series data. Later on we will learn how to work with non-stationary data.

6.1.2 Correlation vs Autocorrelation

In statistics, correlation is a measure of relationship between two variables. In the time series setting, we can think of the current period value and the past period value of a variable as two **separate** variables, and compute correlation between them. Such a correlation, between current and lagged observation of a time series is called **serial correlation** or **autocorrelation**. In general, for a time series, $\{y_t\}$, the autocorrelation is given by:

$$Cor(y_t, y_{t-s}) = \frac{Cov(y_t, y_{t-s})}{\sqrt{\sigma_{y_t}^2 \times \sigma_{y_{t-s}}^2}} \quad (6.1)$$

where $Cov(y_t, y_{t-s}) = E(y_t - \mu_{y_t})(y_{t-s} - \mu_{y_{t-s}})$ and $\sigma_{y_t}^2 = E(y_t - \mu_{y_t})^2$

For a stationary time series, using the three conditions the **Autocorrelation Function (ACF)** denoted by $\rho(s)$ is given by:

$$ACF(s) \text{ or } \rho(s) = \frac{\gamma(s)}{\gamma(0)} \quad (6.2)$$

Non-zero values of the ACF indicates presences of serial correlation in the data. Figure 6.3 shows the ACF for a stationary time series with positive serial correlation. If your data is stationary then the ACF should eventually converge to 0. For a non-stationary data, the ACF function will not decay over time.




Figure 6.4: PACF for a Stationary Time Series

6.1.3 Partial Autocorrelation

Definition 6.2 (Partial Auto Correlation Function (PACF)).

The ACF captures the relationship between the current period value of a time series and all of its past observations. It includes both direct as well as indirect effects of the past observations on the current period value. Often times it is of interest to measure the direct relationship between the current and past observations, **partialing** out all indirect effects. The *partial autocorrelation function (PACF)* for a stationary time series y_t at lag s is the direct correlation between y_t and y_{t-s} , after filtering out the linear influence of $y_{t-1}, \dots, y_{t-s-1}$ on y_t . Figure 6.4 below shows the PACF for a stationary time series where only one lag directly affects the time series in the current period.

6.1.4 Lag operator

A **lag operator** denoted by L allows us to write ARMA models in a more concise way. Applying lag operator once moves the time index by one period; applying it twice moves the time index back by two period; applying it s times moves the index back by s periods.

$$Ly_t = y_{t-1}$$

$$L^2y_t = y_{t-2}$$

$$L^3y_t = y_{t-3}$$

$$\vdots$$

$$L^sy_t = y_{t-s}$$

6.2 Autoregressive (AR) Model

A *stationary* time series $\{x_t\}$ can be modeled as an AR process. In general, an AR(p) model is given by:

$$y_t = \phi_0 + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t \quad (6.3)$$

Here ϕ_i captures the effect of y_{t-i} on y_t . The order of the AR process is not known apriori. It is common to use either AIC or BIC to determine the optimal lag length for an AR process.

Using the Lag operator, we can rewrite the above AR(p) model as follows:

$$\Phi(L)y_t = \phi_0 + \epsilon_t$$

where $\Phi(L)$ is a polynomial of degree p in L :

$$\Phi(L) = 1 - \phi_1 L - \phi_2 L^2 - \dots - \phi_p L^p$$

For example, an AR(1) model can be written as:

$$y_t = \phi_0 + \phi_1 y_{t-1} + \epsilon_t \Rightarrow \Phi(L)y_t = \phi_0 + \epsilon_t$$

where,

$$\Phi(L) = 1 - \phi_1 L$$

Characteristic equation: A characteristic equation is given by:

$$\Phi(L) = 0$$

The roots of this equation play an important role in determining the dynamic behavior of a time series.

6.2.1 Unit root and Stationarity

For a time series to be stationary there should be no **unit root** in its *characteristic equation*. In other words, all roots of the characteristic equation must fall outside the unit circle. Consider the following AR(1) model:

$$\Phi(L)y_t = \phi_0 + \epsilon_t$$

The characteristic equation is given by:

$$\Phi(L) = 1 - \phi_1 L = 0$$

The root that satisfies the above equation is:

$$L^* = \frac{1}{\phi_1}$$

For no unit root to be present, $L^* > |1|$ which implies that $|\phi_1| < 1$.

Typically, for any AR process to be stationary, some restrictions will be imposed on the values of ϕ_i 's, the coefficients of the lagged variables in the model.

6.2.2 Properties of an AR(1) model

A stationary AR(1) model is given by:

$$y_t = \phi_0 + \phi_1 y_{t-1} + \epsilon_t \quad ; \quad \epsilon_t \sim WN(0, \sigma^2) \text{ and } |\phi_1| < 1$$

1. ϕ_1 measures the persistence in data. A larger value indicates shocks to y_t dissipate slowly over time.
2. Stationarity of y_t implies certain restrictions on the AR(1) model.
 - i. Constant long run mean: is the unconditional expectation of y_t :

$$E(y_t) = \mu_y = \frac{\phi_0}{1 - \phi_1}$$

- ii. Constant long run variance: is the unconditional variance of y_t :

$$Var(y_t) = \sigma_y^2 = \frac{\sigma_\epsilon^2}{1 - \phi_1^2}$$

- iii. ACF function:

$$\rho(s) = \phi_1^s$$

6.3 Estimating an AR model

When estimating the AR model we have two alternatives: 1. OLS: biased (but consistent) estimates. Also, later on when we add MA components we cannot use OLS.

2. Maximum Likelihood Estimation (MLE): can be used to estimate AR as well as MA components

6.3.1 Maximum Likelihood Estimation (MLE)

- MLE approach is based on the following idea:

what set of values of our parameters maximize the likelihood of observing our data if the model we have was used to generate this data.

Likelihood function: is a function that gives us the probability of observing our data given a model with some parameters.

6.3.1.1 Likelihood vs Probability

Consider a simple example of tossing a coin. Let X denotes the random variable that is the outcome of this experiment being either heads or tails. Let θ denote the probability of heads which implies $1 - \theta$ is the probability of obtaining tails. Here, θ is our parameter of interest. Suppose we toss the coin 10 times and obtain the following data on X :

$$X = \{H, H, H, H, H, H, T, T, T, T\}$$

Then, the probability of obtaining this sequence of X is given by:

$$Prob(X|\theta) = \theta^6(1 - \theta)^4$$

This is the probability distribution function the variable X . As we change X , we get a different probability for a given value of θ .

Now let us ask a different question. Once we have observed the sequence of heads and tails, let's call it our data which is fixed. Then, what is probability of observing this data, if our probability distribution function is given by the equation above? That gives us the likelihood function:

$$L(\theta) = \text{Prob}(X|\theta) = \theta^6(1 - \theta)^4$$

Note that with fixed X , as we change θ the likelihood of observing this data will change.

This is an important point that distinguishes likelihood function from the probability distribution function. Although both have the same equation, the probability function is a function of the data with the value of the parameter fixed, while the likelihood function is a function of the parameter with the data fixed.

6.3.1.2 Maximum Likelihood Estimation

Now we are in a position to formally define the likelihood function.

Definition 6.3. Let X denotes a random variable with a given probability distribution function denoted by $f(x_i|\theta)$. Let $D = \{x_1, x_2, \dots, x_n\}$ denote a sample realization of X . Then, the likelihood function, denoted by $L(\theta)$ is given by:

$$L(\theta) = f(x_1, x_2, \dots, x_n|\theta)$$

If we further assume that each realization of X is independent of the others, we get:

$$L(\theta) = f(x_1, x_2, \dots, x_n|\theta) = f(x_1|\theta) \times f(x_2|\theta) \times \dots \times f(x_n|\theta)$$

A mathematical simplification is to work with natural logs of the likelihood function, which assuming independently distributed random sample, gives us:

$$\ln L(\theta) = \ln(f(x_1|\theta) \times f(x_2|\theta) \times \dots \times f(x_n|\theta)) = \sum_{i=1}^N \ln(f(x_i|\theta))$$

Definition 6.4. The maximum likelihood estimator, denoted by $\hat{\theta}_{MLE}$, maximizes the log likelihood function:

$$\hat{\theta}_{MLE} \equiv \arg \max_{\theta} \ln L(\theta)$$

Example 6.1. Compute maximum likelihood estimator of μ of an independently distributed random variable that is normally distributed with a mean of μ and a variance of 1:

$$f(y_t|\mu) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(y_t - \mu)^2}$$

Solution: The log likelihood function is given by:

$$\ln L = -T \ln 2\pi - \frac{1}{2} \sum_{t=1}^T (y_t - \mu)^2$$

From the first order condition, we get

$$\frac{\partial \ln L}{\partial \mu} = \sum_{t=1}^T (y_t - \mu) = 0 \Rightarrow \hat{\mu}_{MLE} = \frac{\sum_{t=1}^T y_t}{T}$$

6.3.2 MLE of an AR(p) model

One complication we face in estimating an AR(p) model is that by definition the realizations of the variable are not independent of each other. As a result we cannot simplify the likelihood function by multiplying individual probability density functions to obtain the joint probability density function, i.e.,

$$f(y_1, y_2, \dots, y_T | \theta) \neq f(y_1 | \theta) \times f(y_2 | \theta) \times \dots \times f(y_T | \theta)$$

Furthermore, as the order of AR increases, the joint density function we need to estimate becomes even more complicated.

6.3.3 Forecasting using AR(p) model

Having estimated our AR(p) model with the optimal lag length, we can use the conditional mean to compute the forecast and conditional variance to compute the forecast errors. Consider an AR(1) model:

$$y_t = \phi_0 + \phi_1 y_{t-1} + \epsilon_t$$

Then, the 1-period ahead forecast is given by:

$$f_{t,1} = E(y_{t+1} | \Omega_t) = \phi_0 + \phi_1 y_t$$

Similaryl, the 2-period ahead forecast is given by:

$$f_{t,2} = E(y_{t+2} | \Omega_t) = \phi_0 + \phi_1 E(y_{t+1} | \Omega_t) = \phi_0 + \phi_1 f_{t,1}$$

In general, we can get the following recursive forecast equation for h-period's ahead:

$$f_{t,h} = \phi_0 + \phi_1 f_{t,h-1}$$

Correspondingly, the 1-period ahead forecast error is given by:

$$e_{t,1} = y_{t+1} - f_{t,1} = \epsilon_{t+1}$$

The 2-period ahead forecast error is given by:

$$e_{t,2} = y_{t+2} - f_{t,2} = \epsilon_{t+2} + \phi_1 \epsilon_{t+1}$$

Hence, the h-period ahead forecast is given by:

$$e_{t,h} = \epsilon_{t+h} + \phi_1 \epsilon_{t+h-1}$$

Theorem 6.1. *The h-period ahead forecast converges to the unconditional mean of y_t , i.e.,*

$$\lim_{h \rightarrow \infty} f_{t,h} = \mu_y = \frac{\phi_0}{1 - \phi_1}$$

6.4 Moving Average (MA) Model

A *stationary* time series $\{y_t\}$ can be modeled as an MA(q) process:

$$y_t = \theta_0 + \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q} \quad (6.4)$$

6.5 ARMA(p, q)

An ARMA model simply combines both AR and MA components to model the dynamics of a time series. Formula,

$$y_t = \phi_0 + \phi_1 x_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q} \quad (6.5)$$